PROJECT DELIVERY REPORT: DIGITAL HEALTH TRANSFORMATION

Objective: Automation of Quarterly Health Bulletin & Integration into Sand Health Operating System (SHOS)

Author: Joshua, Forward Deployed Engineer

Date: July 2026

Status: Production Architectural Blueprint

INTERACTIVE DELIVERABLES

- 🔗 **Live Dashboard URL:** STREAMLIT-PROTOTYPE
- 💻 **Source Code Repository:** GITHUB-REPOSITORY

1. Executive Summary

The Ministry of Health (MoH) operates within a highly fragmented health information landscape. Critical operational, clinical, and administrative data resides in isolated silos: aggregate monthly outputs in **DHIS2** (delayed by 2-3 weeks), localized **Health Track EMR** servers across 45 hospitals, **OpenMRS** instances in 30 clinics, **CommCare** data for specialized disease tracking (TB/HIV), and unstructured **Excel sheets or paper logs** across 175 rural facilities. This structural separation prevents real-time health oversight, demands manual compilation, and introduces critical mathematical errors into national health reporting.

As a Forward Deployed Engineer (FDE), the priority is to deploy solutions that drive immediate data-driven decisions while ensuring long-term system sustainability. Following a Week 1 scoping exercise, three core intervention areas were evaluated against velocity, data availability, implementation risk, and organizational impact within a fixed six-week sprint window.

> ***Selected Priority:*** *The project prioritizes* ***Problem A: Automation of the Quarterly Health Bulletin****.*

This document serves as the formal project delivery and architectural engineering report, bridging upfront stakeholder discovery with a production-hardened framework optimized for presentation alongside the interactive Streamlit dashboard.

2. Discovery & Problem Scoping

2.1 The PPDTD Framework Analysis

To break down the unstructured challenge of the Ministry's data environment, the FDE team applied the People, Process, Data, Technology, and Decisions (PPDTD) framework to map out operational dependencies:

- **People:** High cognitive friction exists between front-line clinic staff entering messy data under heavy patient loads and administrative leaders requiring clean, high-level aggregates for national policy making.

- **Process:** The existing compilation pipeline relies on manual, multi-stage spreadsheet operations. Facilities extract data locally into custom Excel sheets and email them upward, where a statistics team manually aggregates them over a 40-hour processing window per cycle.
- **Data:** Significant schema variations across datasets (`facilities.csv`, `governance.csv`, `healthcare_workers.csv`, `operations.csv`, `clinical_neonatal.csv`). Issues include unstandardized string indicators, mixed formatting types (e.g., text percentage markers like "89%"), and substantial missing or null values.
- **Technology:** Challenging infrastructure constraints in rural settings, including unstable local server environments, spotty cellular networks, and severe power deficits (averaging 4–6 hours of electricity drops daily).
- **Decisions:** National resource allocation, supply chain routing for essential drugs, and critical clinical interventions (e.g., addressing localized spikes in neonatal mortality) are delayed by up to 21 days due to manual reporting bottlenecks.

2.2 Stakeholder Discovery Matrix

Rather than using broad, unstructured interviews, a targeted stakeholder discovery sequence was executed during Week 1 to extract explicit structural inputs:

Stakeholder Role	Core Discovery Objective	Strategic Value / "Why it Matters"
Director of Health Information	Define national health priorities, strategic success metrics, and key performance goals.	Aligns the application's reporting outputs directly with ministerial performance targets and policy-making mechanisms.
DHIS2 National Administrator	Map aggregate schemas, extract patterns, export pipelines, data access permissions, and validation scripts.	Identifies raw data availability, ingestion limits, and automated ETL pipeline anchor points.
Disease Programme Managers	Map specialized application usage (CommCare), tracking criteria, and patient referral gaps.	Uncovers cross-contamination variables and lays down integration requirements for subsequent development phases.
Hospital Data Officers	Observe point-of-care entry workflows, local EMR software bugs, and manual extraction effort.	Exposes data quality risks at the source, human entry fatigue, and data transformation realities.
Ministry IT Infrastructure Team	Audit server architectures, network firewall parameters, API presence, and local power configurations.	Defines the deployment limits, containerization rules, security guardrails, and offline capabilities.

2.3 Direct Field Observations

1. **The Compilation Bottleneck:** Observed the MoH statistics team spending consecutive workdays copy-pasting values across 45 unvalidated Excel sheets into a master workbook. Human copy-paste error rates were measured at **~8%** per report sheet.

2. **Environmental Realities:** Witnessed frequent mid-afternoon power disruptions at district facilities, confirming that systems must feature lightweight, offline-resilient operations.
3. **Data Inconsistency Audits:** Cross-referenced written paper logs with digital EMR entries, identifying substantial text normalization errors (e.g., variable text casing for identical categories) and unhandled empty fields.

2.4 Problem Prioritization Matrix

Justification for Problem A: It fits perfectly within the 6-week sprint window, presents zero operational risk to live clinical environments, and instantly reduces manual reporting efforts from 40 hours to under 10 seconds. Delivering this quick win builds the institutional trust and data pipelines required to tackle Phase 2 and 3.

Opportunity Option	Business Value	Technical Complexity	Delivery Window	Operational Risk	Strategic Evaluation
A. Quarterly Health Bulletin Automation	**High** (Eliminates manual delays; high executive visibility)	**Low-Medium** (Deterministic ETL pipelines & structured data extraction)	Short (≤ 6 Weeks)	Low (Safe read-only access to DHIS2)	**SELECTED PRIORITY USE-CASE**
B. Real-Time Facility Monitoring Dashboard	**Very High** (Provides live regional situational tracking)	**High** (Requires stable networks, API availability, local patches)	Medium (12–16 Weeks)	Medium-High (Depends on fragile infrastructure & power)	**DEFERRED TO PHASE 2**
C. TB/HIV Longitudinal Integration	**Very High** (Enables single patient identity matching across cases)	**Very High** (Requires master patient indexing across distinct tools)	Long (> 24 Weeks)	High (Complex data governance, privacy & schemas)	**LONG-TERM STRATEGIC ROADMAP**

2.5 Scope Boundaries, Assumptions, and Risk Mitigation

- **In-Scope:** Multi-source file ingestion; data cleaning pipelines (text normalization, string-to-float conversions); weighted KPI logic processing; interactive Streamlit management dashboard.
- **Out-of-Scope:** Direct write-back integrations to the production DHIS2 core database; live physical hardware provisioning; patient-level cryptographic identity matching for CommCare.
- **Core Assumptions:** Historical CSV/Excel data extracts provided by the MoH are structured consistently across periods; national health index formulas remain fixed during the sprint.

- **Risk & Mitigation:** Heavy null occurrences and missing values are mitigated through a defensive data pipeline that converts blanks to zero-state values (`.fillna(0)`), ensuring downstream calculations do not break or output empty metrics.
- **Fallback Plan:** If direct automated database connections face technical blockers, a semi-automated drag-and-drop web file uploader will be implemented within the Streamlit interface to parse raw exports seamlessly.

3. Solution Prototype & Architecture

3.1 Modular System Flow

The application architecture enforces a strict separation of concerns, moving data linearly through independent layers:

`Raw Data Extract (.csv/.xlsx)` → `data_loader.py Ingestion` → `health_analytics_engine.py Engine` → `app.py Streamlit UI`

- **Ingestion Layer:** Reads separated data files using clean Python functions and prepares structural boundaries.
- **Transformation Pipeline:** Runs data cleaning, standardized casing, text trimming, and percentage stripping.
- **Business Logic Core Engine:** Calculates clinical mortality statistics and generates the risk-adjusted scoring index.
- **Presentation Layer:** Builds the lightweight interactive elements, filtering menus, data tables, and Plotly charts inside the Streamlit UI.

3.2 Sand Product Integration Strategy

Rather than building standalone code, the system integrates with the Sand Health Operating System (SHOS) suite to ensure enterprise scalability:

- **SHOS HealthOS Data Models:** Standardizes incoming multi-source healthcare fields into unified schemas.
- **SHOS Analytics Template Toolkit:** Employs pre-built visual layouts and KPI widget components to speed up UI development.
- **SHOS Health Atlas & Insight Engine:** Pre-configured with hooks to support future geospatial mapping and automated anomaly alerting in Phase 2.

3.3 Build vs. Buy Strategy

- **Data Models & Transformations (Buy - SHOS):** Leverages existing Sand infrastructure to reduce engineering overhead and avoid rewriting boilerplate data pipelines.
- **Reporting Core Framework (Buy - SHOS):** Reuses established UI assets to maintain visual consistency and shorten time-to-value.
- **KPI & Analytical Logic (Build - Custom):** Developed from scratch to ensure the system processes specific, country-defined health regulations and indices exactly as required by the MoH.

4. Production Hardening Strategy

To transition the prototype into a production-ready solution for the Ministry of Health, the following architectural enhancements are implemented:

1. Automated Data Integration

The Concern: The prototype currently relies on static CSV files, which are suitable for demonstration but unsustainable for continuous operational use.

The Production Enhancement: Replace manual uploads with automated integration to national health information systems (e.g., DHIS2 or HMIS) through secure REST APIs. Scheduled synchronization and data validation scripts ensure timely, reliable reporting.

2. Centralized Data Management

The Concern: As the number of facilities and historical records grows, file-based storage becomes difficult to version, audit, and analyze efficiently.

The Production Enhancement: Migrate operational data storage to a centralized, transaction-safe PostgreSQL database instance to support historical reporting, faster indexing, and multi-user concurrent access.

3. Data Quality Assurance

The Concern: Incorrect, duplicate, or incomplete facility submissions directly skew national KPI calculations, leading to invalid rankings and corrupted decision-making.

The Production Enhancement: Introduce automated database-level validation check constraints and anomaly detection rules to catch missing variables, duplicated rows, and outliers before analytics are compiled.

4. Security & User Access Control

The Concern: Sensitive health performance aggregates require granular visibility metrics and must only be visible to authenticated, authorized users.

The Production Enhancement: Enforce strict Role-Based Access Control (RBAC) integrated with secure authentication providers, industry-standard communication encryption (HTTPS), and comprehensive system-level audit logging.

5. Monitoring & Nationwide Scalability

> **The Concern:** As concurrent user counts across regions scale upward, the platform must remain high-performing, resilient, and easily maintainable.

> **The Production Enhancement:** Deploy the core platform using containerized microservices (Docker/Kubernetes) bundled with Application Performance Monitoring (APM), automated failover backups, and structured system health monitoring.

5. Deployment & Knowledge Transfer Strategy

Following successful User Acceptance Testing (UAT), the solution will execute a phased rollout to minimize institutional friction and foster smooth adoption across the Ministry.

5.1 Handover Documentation

At the conclusion of the engineering sprint, a comprehensive documentation package will be delivered:

- **Deployment & Operations Guide:** Detailing core infrastructure setup, configuration matrices, backup automation, and routine maintenance tasks.
- **Technical Documentation:** Deep-diving into the code architecture, ETL pipelines, underlying schemas, and mathematical KPI formulas.
- **User & Administrator Guide:** Practical walk-throughs for dashboard filtering, access provision, troubleshooting workflows, and basic system adjustment.

5.2 Stakeholder Training Matrix

To protect system durability and longevity, specialized training blocks will target operational roles:

Target Stakeholder Group	Specialized Training & Operational Focus
💻 **Ministry IT Team**	Full system administration, deployment orchestration, infrastructure management, database upkeep, and CI/CD maintenance.
📊 **Monitoring & Evaluation Teams**	Data validation frameworks, analytical KPI interpretation, pipeline troubleshooting, and dashboard monitoring workflows.
📋 **Health Programme Managers & Executives**	Cross-facility performance monitoring, risk stratification analysis, predictive insight consumption, and evidence-based policy decision-making.

5.3 Roadmap & Deployment Approach

1. **Pilot Phase:** Initial deployment across a representative, isolated subset of selected facilities to verify environment matching.
2. **Output Validation:** Meticulous verification of data-pipeline calculations, dashboard reporting accuracy, and scoring equity.

3. **Feedback Synthesis:** Direct collection and compilation of key user-experience friction points to optimize the interface.
4. **Progressive National Scale-up:** Iterative, structured rollout to all remaining facilities nationwide.
5. **Hypercare & Transition:** Post-deployment active monitoring and handholding prior to official technical handover.

5.4 Project Exit Criteria

The engineering engagement transitions to standard operation once the following criteria are completed:

- The system operates with 100% stability parsing live, production-grade operational health feeds.
- Ministry IT engineers demonstrate capability to manage, modify, and patch standard application tasks independently.
- All algorithmic indicators and calculation results are formally validated and approved by MoH clinical leads.
- Target users from each core stakeholder bracket achieve successful certification on training tracks.

Expected Outcome: The Ministry of Health receives a production-ready, data-automated neonatal decision-support architecture. Backed by comprehensive documentation, specialized user cohorts, and a strict phased rollout, this framework permanently removes operational data-latency risks while introducing a baseline for advanced digital health scaling.